

Enterprise Intelligence & Investment Research Platform

Project Context Handbook (Plain-English + Technical)

Prepared for: Akash

Purpose: Give a full-context explanation of the project so a developer, PM, or multi-agent coding system can understand what it is, why it exists, what must be built, and how the pieces fit together.

1) Executive Summary

This project is an **enterprise intelligence and investment research platform**.

In simple language, it is software that lets a user enter a:

- company
- stock ticker
- person
- investor
- fund
- government agency
- bill
- regulation
- docket
- contract
- grant
- court case
- PAC
- nonprofit
- portfolio

...and the platform will:

1. resolve the correct entity,
2. collect evidence from public records, financial data, market data, and licensed data,
3. build a relationship graph,
4. analyze financial, legal, regulatory, procurement, lobbying, and market implications,
5. generate evidence-backed reports,
6. let analysts edit and improve those reports,
7. verify claims and citations again after edits,
8. route sensitive outputs through review and approval,
9. export final deliverables,

10. keep monitoring the entity or portfolio for future changes.

This is **not** just a stock screener, **not** just a public-record search engine, and **not** just an LLM report generator.

It is a combination of:

- investment research workstation,
 - public-record intelligence system,
 - evidence vault,
 - entity resolution engine,
 - relationship graph platform,
 - report editor and review workspace,
 - portfolio and watchlist monitoring system,
 - enterprise collaboration and governance layer.
-

2) What the product is really trying to do

The core idea is:

Turn messy external data into trustworthy, cited intelligence.

Most research tools stop at one of these layers: - they show raw records, - or they provide analysis without evidence, - or they generate text without a review workflow, - or they focus only on market data and ignore lobbying / regulation / procurement / litigation.

This platform combines all of them.

Core promise

A user should be able to ask: - “Analyze this stock.” - “Who is connected to this company?” - “Does this company depend on government contracts?” - “What lobbying or regulatory exposure matters here?” - “What litigation or sanctions risk exists?” - “Can you produce a cited investment memo?” - “Can you keep watching this portfolio for new developments?”

...and get a structured, evidence-backed answer.

3) Who this product is for

The platform is designed for **enterprise and professional users**, not casual retail users.

Supported roles

- **Organization owner** - tenant admin, billing, source policy, global controls
- **Workspace admin** - manages workspace, invites users, approves exports
- **Senior analyst** - leads investigations, approves reports, reviews claims, handles merges
- **Analyst** - researches entities, runs workflows, drafts reports, adds notes/evidence
- **Investor user** - runs stock, valuation, and portfolio workflows
- **Compliance reviewer** - checks sensitive claims, source restrictions, export risk
- **Guest reviewer** - limited review/comment access

User types in business terms

- equity researchers
 - buy-side or long-only analysts
 - event-driven / catalyst-driven researchers
 - due diligence teams
 - policy / regulatory researchers
 - procurement and government-exposure analysts
 - legal-risk analysts
 - compliance teams
 - portfolio managers / investment committees
 - enterprise strategy / diligence teams
-

4) Product philosophy and non-negotiable principles

The uploaded specification makes a few principles very clear. These are not optional.

4.1 Evidence first

Every relationship, risk flag, valuation point, and report claim must trace back to evidence.

That means: - no unsupported claims, - no mystery facts, - no “LLM said so” outputs, - every important conclusion must be source-backed.

4.2 Official source first

Use official APIs and bulk datasets before scraping. Scraping should be limited to lawful, permitted, public pages when no official API exists.

4.3 Human review for high-risk outputs

AI can draft, summarize, or assist. But sensitive conclusions, investment-facing outputs, export-ready reports, and potentially damaging claims must go through review.

4.4 No unsupported allegations

The system should explain what the data shows, not invent accusations. Uncertainty must be labeled clearly.

4.5 Investor decision support, not auto-trading

This platform supports research and analysis. It should **not** auto-execute trades or present guaranteed outcomes.

4.6 Multi-user enterprise governance

Actions must be: - permissioned, - logged, - reviewable, - auditable, - versioned, - reversible where appropriate.

4.7 Cost-aware architecture

The platform should not shove everything into the most expensive database. It should store and process data according to workload and value.

5) Scope: Phase 1 vs Phase 2

5.1 Phase 1

Phase 1 is a **U.S.-focused MVP**, but “MVP” here means **serious first product**, not toy prototype.

Phase 1 must deliver: - enterprise workspaces and governance, - U.S. public-record ingestion, - evidence vault, - entity resolution, - graph intelligence, - investor intelligence, - report generation and editing, - Claude finance skills gateway, - watchlists, alerts, and portfolio monitoring.

Main U.S. source families in Phase 1

- SEC EDGAR
- FEC / OpenFEC
- LDA.gov
- FARA
- Congress.gov
- GovInfo
- Federal Register

- Regulations.gov
- eCFR
- RegInfo/OIRA
- USAspending
- SAM.gov
- IRS Form 990
- CourtListener / RECAP
- OFAC / OpenSanctions
- OpenCorporates
- GLEIF
- market-data providers
- licensed research / market sources where allowed

5.2 Phase 2

Phase 2 is **global expansion**.

It reuses the same: - ontology, - evidence model, - entity resolution model, - graph model, - report system, - review system.

It then adds: - U.K. sources - EU sources - Canada - Australia - Brazil - ICIJ / offshore / investigative sources - global sanctions - cross-border entity resolution - multilingual names / transliteration / currency handling / country-specific identifiers

5.3 Recommended execution approach

Recommendation

Build **Phase 1 first**, but keep the architecture **Phase-2-ready**.

That means: - do not hard-code everything to U.S.-only assumptions, - do include jurisdiction-aware identifiers and source contracts where practical, - but do **not** attempt full global rollout in the first implementation wave.

6) What the platform must do end-to-end

A complete target flow looks like this:

1. User searches for a company, ticker, person, fund, agency, bill, docket, contract, or portfolio.
2. System resolves the entity with identifiers, aliases, and confidence.
3. System gathers evidence from operational stores, raw source data, search indexes, graph edges, market data, and licensed sources.
4. User chooses a workflow:
 - Analyze Stock
 - Analyze Company

- Investor Memo
 - Government Exposure
 - Related Parties
 - Due Diligence Pack
 - Proposal / Pitch
 - Portfolio Report
 - Custom Investigation
5. Platform builds an **Evidence Bundle**.
 6. Platform runs financial, public-record, graph, and optional Claude finance skill workflows.
 7. System generates a draft report with sections, claims, assumptions, charts, tables, and evidence references.
 8. Draft enters **Review Workspace**.
 9. Analysts edit or enhance sections.
 10. Every factual edit triggers claim re-verification.
 11. Reviewer approves or rejects sections and claims.
 12. Final output is exported to PDF / Word / PPT / Excel / Markdown / HTML / JSON / evidence appendix.
 13. User can add entity or portfolio to watchlists for future alerts.
-

7) Major product modules

Below is the best plain-language explanation of the major modules.

7.1 Enterprise Workspace and Identity

This is the enterprise shell of the product.

It includes: - organizations - workspaces - users - roles - permissions - projects - cases - watchlists - collaboration - audit logs - SSO-ready architecture - SCIM-ready provisioning hooks - MFA-ready design - tenant isolation

This is what makes the product usable by teams instead of a single person.

7.2 Source Registry and Connector Framework

This is the ingestion control layer.

It keeps track of: - what sources exist, - how to authenticate, - how often they refresh, - what their terms/compliance constraints are, - what connector last ran, - what failed, - what needs retry, - what has changed.

Every source should have a reusable **source contract** defining: - auth, - rate limits, - refresh cadence, - compliance notes, - raw storage requirements, - normalized outputs.

7.3 Evidence Vault

This is the most important trust layer.

The Evidence Vault stores: - raw API responses, - raw JSON/XML/CSV/HTML/PDF/text, - document hashes, - retrieval metadata, - source-native IDs, - EvidenceRefs, - snippets/excerpts/page spans, - immutable raw artifacts.

Why it matters: If a report says something serious, the platform must be able to say: - where it came from, - when it was fetched, - what raw source it came from, - and what exact excerpt supports it.

7.4 U.S. Source Connectors

These are the data-ingestion adapters for official and public sources.

Examples: - SEC EDGAR - FEC - LDA.gov - FARA - Congress.gov - GovInfo - Federal Register - Regulations.gov - eCFR - RegInfo/OIRA - USAspending - SAM.gov - IRS 990 - CourtListener - OpenSanctions / OFAC - OpenCorporates - GLEIF

Each connector should: 1. discover records, 2. fetch records, 3. store raw first, 4. parse fields, 5. normalize entities, 6. create EvidenceRefs, 7. create candidate relationships, 8. trigger entity resolution, 9. index useful text, 10. emit data-quality checks and metrics.

7.5 Entity Resolution

This is how the system decides that multiple records refer to the same real-world thing.

Deterministic identifiers

Use exact IDs when available, such as: - CIK - LEI - EIN - UEI - CAGE - FEC IDs - LDA IDs - FARA IDs - court IDs - award IDs - RIN - docket IDs - OpenCorporates company numbers

Probabilistic / fuzzy signals

Use these when exact IDs are missing: - name similarity - alias similarity - address overlap - officer overlap - parent/sub relationship overlap - website/domain overlap if public - co-occurrence in filings - graph-neighborhood similarity - temporal plausibility

Merge rules

- high-confidence deterministic matches can auto-merge,
 - medium-confidence matches go to review,
 - low-confidence matches create possible-match edges,
 - conflicting IDs block auto-merge,
 - every merge must be reversible.
-

7.6 Intelligence Graph

This is the relationship engine of the product.

It turns entity records into an evidence-backed graph of connections such as:
- ownership, - investment, - officers, - boards, - lobbying, - political money, - contract awards, - grants, - rulemaking, - public comments, - OIRA meetings, - litigation, - sanctions, - related-party links.

The graph must support: - graph expansion, - pathfinding, - related-party discovery, - timelines, - evidence-backed edge inspection, - graph visualization.

This is one of the main features that makes the product useful for investigations and due diligence.

7.7 Investor Intelligence

This is the financial analysis layer.

It supports: - stock analysis, - company analysis, - investor memo generation, - valuation, - DCF, - comps, - technical analysis, - catalysts, - ownership, - earnings analysis, - risk factors, - monitoring plans.

Example analysis modes

- day trader
- swing trader
- long-term investor
- value investor
- growth investor
- event-driven
- activist
- short thesis
- credit
- private equity
- government-contracting
- regulatory risk

Required outputs

- company overview
 - stock snapshot
 - price/volume trend
 - liquidity and volatility
 - fundamentals
 - valuation summary
 - DCF assumptions and range
 - peer comps
 - technicals
 - earnings
 - insiders
 - institutional ownership
 - government exposure
 - lobbying exposure
 - regulatory exposure
 - procurement exposure
 - litigation / enforcement risk
 - sanctions / exclusions
 - related parties
 - bull / base / bear case
 - catalysts
 - risks
 - monitoring plan
 - evidence appendix
-

7.8 Finance Skills Gateway (Claude + internal fallback)

The project should support Claude's financial-services skills and agents as an optional but first-class layer.

Examples of workflows mentioned in the uploaded specs: - comparable-company analysis - discounted cash flow models - earnings analysis - company teaser / one-pager - initiating coverage - due-diligence pack - valuation review - pitch / proposal support - market research

Important design rule

Do **not** hard-code the product to one provider.

Instead build a **provider-agnostic FinanceSkillsGateway** that can route to:
- Anthropic finance workflows when enabled - internal fallback skills/templates when needed - mock/test providers for development

Why this matters

- keeps the product portable,
 - lets the team operate when vendor access is unavailable,
 - makes testing easier,
 - makes auditing easier,
 - keeps the rest of the app independent from one model vendor.
-

7.9 Report Generation and Review

This is not a simple “generate report” button. It is a full **draft** → **edit** → **verify** → **review** → **export** workflow.

Core report capabilities

- report templates
- evidence bundles
- claims
- claim-evidence links
- citation verifier
- contradiction scanning
- rich editor
- tracked changes / redlines
- comments
- version history
- report enhancement actions
- review tasks
- final approval
- exports

Report types in Phase 1

- stock analysis
- investor memo
- initiating coverage
- earnings update
- DCF valuation
- company intelligence
- government exposure
- lobbying influence
- procurement risk
- regulatory exposure
- litigation risk
- related-party report
- portfolio review

- proposal / pitch
- vendor / counterparty due diligence

Critical rule

If a user edits factual content and removes evidence, that section must be marked as needing evidence again.

7.10 Alerts and Monitoring

This keeps the research alive after the report is done.

The platform should allow users to watch: - tickers - companies - people - funds - agencies - bills - dockets - contractors - lobbyists - keywords - portfolios - cases

Example triggers

- new SEC filing
- new Form 4
- new 13D/G
- new 13F
- new Form D
- new lobbying report
- bill action
- Federal Register item
- public comment
- OIRA meeting
- contract award
- SAM opportunity
- court case
- sanctions hit
- nonprofit grant
- related-party event
- price movement
- volume spike
- earnings event
- news spike
- valuation threshold

Delivery

- in-app
- email
- Slack
- Teams
- webhook

- API
 - mobile later
-

7.11 Portfolios

The platform also needs portfolio workflows.

Portfolio capabilities include: - CSV upload - manual position entry - mapping positions to entities - watchlists - unrealized P/L - sector exposure - factor exposure - regulatory exposure - government-contract exposure - litigation/sanctions exposure - earnings calendar - catalyst calendar - portfolio memo - thesis tracker

This lets the platform move beyond single-name analysis into portfolio-level monitoring.

7.12 Admin, Compliance, and Governance

This is the enterprise-control module.

It must handle: - source credentials - source policies - restricted fields - export policies - retention policies - legal hold hooks - audit logs - security settings - model-provider settings - billing hooks - source restrictions - high-risk claim approvals

This matters because the platform deals with: - sensitive public records, - regulated financial contexts, - potentially damaging claims, - tenant-scoped uploads, - licensed data.

8) Key workflows a non-finance person should understand

8.1 Analyze Stock

User enters a ticker or company. System combines: - market data, - SEC filings, - valuation, - technicals, - related parties, - government exposure, - litigation/regulation/lobbying/procurement insights.

Output: a structured investment research pack.

8.2 Company Intelligence Report

User enters a company. System produces a broad intelligence profile: - identity, - ownership, - officers, - investors, - contracts, - lobbying, - regulations, - lawsuits, - sanctions, - related parties, - timeline.

8.3 Related-Party Investigation

User enters a person, company, or fund. System explores the graph and shows who else matters, why, and what evidence supports it.

8.4 Government Exposure Report

The platform identifies: - contract/grant exposure, - agencies involved, - lobbying overlap, - regulations or dockets, - procurement concentration, - investor implications.

8.5 Portfolio Monitoring

User uploads a portfolio. System tracks events across all names and sends alerts / briefs.

8.6 Report Editing and Enhancement

AI can draft and improve parts of the report, but: - the human remains in control, - claims must remain cited, - high-risk outputs must be reviewed.

9) Architecture explained simply

The specification strongly recommends **not** building this as one giant database or one giant RAG app.

The core architecture idea

Use the right storage engine for the right job.

Recommended logical layers

9.1 Operational OLTP Use **Postgres** for: - users - organizations - workspaces - entities - identifiers - claims - reports - jobs - permissions - connector state

9.2 Raw Evidence Lake Use **object storage** for: - raw API responses - JSON - XML - CSV - PDFs - HTML - filings - transcripts - attachments - generated artifacts

This should use immutable paths and content hashes.

9.3 Lakehouse Use **Parquet + Iceberg** (preferred) or **Delta Lake** for: - parsed source tables - append-only historical tables - bronze/silver/gold layers - large analytical history

9.4 Warehouse / OLAP Use **ClickHouse** first for event-heavy analytics or dashboards. Offer **BigQuery** / **Snowflake** later if enterprise customers require it.

9.5 Search Use **OpenSearch** for: - full-text retrieval - document search - snippets - hybrid retrieval - later vector/hybrid support

9.6 Graph Start with **Postgres edge tables**. Only move to **Neo4j** / **ArangoDB** / **Neptune** later when graph queries actually justify it.

Best current recommendation

For Phase 1: - FastAPI backend - Next.js frontend - Postgres - object storage - Redis / queue - OpenSearch - Parquet export - DuckDB for low-cost analytics - Postgres graph first

This gives the best balance between speed, cost, and maintainability.

10) Why this architecture is smart

The docs are very deliberate about cost.

What to avoid

- putting every raw record into an expensive warehouse,
- embedding every document chunk automatically,
- polling paid APIs too often,
- duplicating raw documents across systems,
- regenerating full reports when only one section changed.

What to do instead

- store raw evidence cheaply in object storage,
- normalize only what is needed,
- cache expensive API calls,
- use hot/warm/cold data strategies,
- create report sections incrementally,
- use budget-aware report generation,
- only add expensive infrastructure when scale justifies it.

This is one of the strongest parts of the project design.

11) Data model: the most important objects

Below are the concepts you must understand to reason about the system.

11.1 Entity

A company, person, fund, agency, bill, regulation, court case, PAC, nonprofit, security, document, or report.

Key fields: - type - canonical_name - aliases - identifiers - jurisdictions - description - confidence

11.2 Relationship

A graph edge connecting two entities.

Examples: - OWNS - INVESTED_IN - OFFICER_OF - BOARD_MEMBER_OF - LOBBIED_FOR - LOBBIED_ON - DONATED_TO - RECEIVED_CONTRACT - RECEIVED_GRANT - PARTY_TO_CASE - SANCTIONED_BY - EXCLUDED_BY - SAME_AS - MENTIONED_IN - RELATED_TO

Every relationship should have: - evidence refs - source record IDs - confidence - dates - amount where relevant - review status

11.3 EvidenceRef

The traceability object.

This links a fact or claim back to: - source name - source record ID - raw document ID - source URL - document type - publication date - retrieved date - excerpt / field path - page number or character range - hash

11.4 Claim

A factual or analytical assertion used in a report.

It must include: - claim text - type - confidence label - evidence refs - contradictions - assumptions - review status - author / reviewer

11.5 SourceRun

One ingestion job for one source, including status, timestamps, metrics, and errors.

11.6 EvidenceBundle

A curated package of: - source records, - raw docs, - normalized facts, - graph paths, - financial metrics, - market events, - assumptions, - quality warnings

used to generate a specific report.

11.7 SkillRun

A logged invocation of a finance skill or agent, including: - provider - skill name
- version - artifacts - prompt hash - evidence refs - cost estimate - actual cost -
status - review state

12) Core APIs the product needs

At a high level, the product will need APIs for:

Search and entities

- global search
- entity profile
- relationships
- graph expansion / path
- timeline
- evidence

Finance

- analyze stock
- DCF
- comps
- technicals
- earnings summary
- portfolio analysis

Sources

- source list
- source run
- source-run status/errors
- source records

Skills

- skill list
- run skill
- approve / reject skill run
- retrieve artifacts

Reports

- create report
- generate draft

- enhance section
- verify claims
- approve/reject claims
- export
- versions / compare

Alerts / watchlists

- create watchlist
- add watchlist items
- create alert
- pause / resume
- list alert events

13) What “good implementation” means for this project

Because this is a complicated platform, a shallow implementation will fail.

A good implementation must be:

13.1 Code-wise, not docs-only

This project cannot be “implemented” with: - markdown plans only, - folder scaffolds only, - fake placeholder services only, - pseudo-APIs only.

It needs: - real services, - real models, - real migrations, - real workers, - real connector code, - tests, - fixtures, - UI flows, - persistence, - auth/permissions, - evidence traceability.

13.2 Evidence-preserving

If evidence is weak or missing, the system must surface that honestly.

13.3 Reviewable

All important outputs should be editable, reviewable, and auditable.

13.4 Cost-aware

Do not build the most expensive version first.

13.5 Extensible

Phase 2 must be addable without re-architecting everything.

14) What a non-finance PM / founder should watch out for

If you do not come from finance or public-record intelligence, these are the biggest project risks:

14.1 Overbuilding before evidence workflows are stable

If the evidence vault and entity model are weak, the rest will look impressive but won't be trustworthy.

14.2 Treating LLM output as final truth

LLM output must stay draft-quality until verified and reviewed.

14.3 Building only a “stock report generator”

That would miss most of the product's real value: - public records, - graph relationships, - government exposure, - litigation/regulation/procurement context, - collaboration/review.

14.4 Putting all data into one expensive storage tier

This will blow up cost fast.

14.5 Weak entity resolution

If the system merges the wrong entities, everything downstream gets corrupted.

14.6 Poor connector discipline

Every source needs: - raw persistence, - tests, - contracts, - compliance notes, - retries/checkpoints, - quality checks.

14.7 Missing review and export controls

This is enterprise software, not a consumer chatbot.

15) Recommended build order

Based on the uploaded specs and prior project analysis, this is the recommended order.

Step 1 - Foundation

- monorepo
- services
- auth

- Postgres
- object storage abstraction
- job runner
- source contract template
- AGENTS rules
- memory/context files

Step 2 - Enterprise core

- organizations
- workspaces
- users
- roles
- permissions
- audit logs
- projects/cases

Step 3 - Evidence platform

- source registry
- source runs
- raw docs
- EvidenceRef
- document viewer

Step 4 - P0 connectors

- SEC
- FEC
- LDA
- Congress.gov
- Federal Register
- USAspending
- OpenCorporates
- GLEIF
- OpenSanctions / OFAC

Step 5 - Entity resolution + ontology

- entity tables
- identifiers
- relationships
- merge/unmerge
- review queue

Step 6 - Graph intelligence

- pathfinding
- related-party discovery
- timeline UI
- graph UI

Step 7 - Investor intelligence

- market data abstraction
- financial statements
- DCF
- comps
- technicals
- stock dashboard

Step 8 - Finance skills gateway

- provider abstraction
- Anthropic adapter
- internal fallback
- skill registry
- artifact storage

Step 9 - Report engine + review workspace

- evidence bundles
- claims
- claim verifier
- editor
- comments
- redlines
- versioning
- exports

Step 10 - Alerts and portfolios

- watchlists
- alert rules
- delivery
- portfolio import
- exposure analytics

Step 11 - P1 connector wave

- FARA
- GovInfo
- Regulations.gov

- eCFR
- RegInfo/OIRA
- SAM.gov
- IRS 990
- CourtListener / RECAP

Step 12 - Enterprise hardening

- OpenSearch
- compliance controls
- SSO/SCIM hooks
- field-level permissions
- export approvals
- observability
- tests
- load tests
- docs

16) Suggested 15 implementation issues (high-level)

Below is a short-form summary of the 15-issue plan we previously derived for LLM / multi-agent execution.

1. Platform foundation, monorepo, memory system, local env
2. Enterprise identity, organizations, workspaces, roles, audit logging
3. Source registry, credentials, source contracts, connector run framework
4. Evidence vault, raw docs, hashing, EvidenceRefs, document viewer
5. P0 connectors wave 1
6. P1 connectors wave 2
7. Canonical ontology, entity model, identifier system, entity resolution
8. Search, entity profiles, timelines, discovery APIs
9. Intelligence graph, pathfinding, related-party discovery, graph UI
10. Investor intelligence: market data, valuation, DCF, comps, technicals
11. Finance skills gateway, Anthropic adapters, internal fallback skills

12. Report engine, evidence bundles, claims, citation verifier, templates
13. Review workspace, editing, enhancement, redlines, versioning, exports
14. Watchlists, alerts, portfolio import, monitoring jobs
15. Search infra, compliance, security hardening, cost controls, testing, docs

Very important rule for any LLM / multi-agent implementation

The project must be implemented **code-wise**, not as markdown-only planning artifacts.

That means: - real code, - real schema, - real APIs, - real workers, - real tests, - real UI, - real configs, - real migrations, - real review flows.

17) Plain-English glossary for someone new to finance

DCF

Discounted Cash Flow. A valuation model that estimates what a business is worth today based on future cash generation.

Comps

Comparable-company analysis. Compare a company to similar companies using valuation multiples.

10-K / 10-Q / 8-K

SEC filing types: - 10-K = annual report - 10-Q = quarterly report - 8-K = important event/update

13F / 13D / 13G

Ownership-related SEC filings.

Form 4

Insider transaction filing.

Catalyst

An event that could move the stock or materially change the thesis.

Lobbying exposure

Whether the company or related firms are active in lobbying around relevant issues/agencies.

Government exposure

How much the company depends on contracts, grants, agencies, or regulatory processes.

Related party

A connected entity or person that could affect risk, influence, control, or due diligence.

EvidenceRef

A structured reference to the specific raw evidence supporting a fact or claim.

Watchlist

A saved monitoring target that triggers alerts when new relevant events happen.

Investigation Workspace

A collaborative area where analysts research entities, collect evidence, build claims, and draft reports.

18) Recommended final positioning of the product

If you need one clean sentence to explain the project internally or externally, use this:

Enterprise Intelligence & Investment Research Platform is a multi-user, evidence-first system that combines public-record intelligence, investment research, graph analysis, report generation, review workflows, and monitoring into one enterprise product.

A more detailed version:

It helps analysts, investors, and enterprise teams investigate companies, people, funds, agencies, regulations, contracts, litigation, sanctions, and portfolios by collecting public and financial evidence, resolving identities, mapping relationships, generating cited reports, and monitoring future developments.

19) Recommended decision summary

Recommended execution decision

- **Build Phase 1 now**
 - **Keep Phase 2-ready architecture**
 - **Prioritize evidence, connectors, entity resolution, graph, and review workflow over flashy AI output**
 - **Use provider-agnostic finance skills integration**
 - **Treat all AI-generated finance outputs as draft until reviewed**
 - **Prefer official APIs and bulk data over scraping**
 - **Keep raw evidence cheap, durable, and immutable**
 - **Add expensive analytics or graph infra only when scale justifies it**
-

20) External references and official docs to keep handy

Official/public source docs

- SEC EDGAR APIs: <https://www.sec.gov/search-filings/edgar-application-programming-interfaces>
- OpenFEC: <https://api.open.fec.gov/developers/>
- LDA.gov API: <https://lda.senate.gov/api/>
- Congress.gov API: <https://api.congress.gov/>
- Federal Register API: <https://www.federalregister.gov/developers/documentation/api/v1>
- USAspending API: <https://api.usaspending.gov/>
- GLEIF API: <https://www.gleif.org/en/lei-data/gleif-api>
- OpenCorporates API: <https://api.opencorporates.com/documentation/API-Reference>
- CourtListener API overview: <https://wiki.free.law/c/courtlistener/help/api/rest/v4/overview>
- OpenSanctions repo: <https://github.com/opensanctions/opensanctions>

Architecture / platform docs

- Apache Iceberg: <https://iceberg.apache.org/>
- Delta Lake: <https://delta.io/>
- OpenSearch Docs: <https://docs.opensearch.org/latest/>
- ClickHouse Docs: <https://clickhouse.com/docs>

Finance-skill / AI references

- Anthropic financial-services repository: <https://github.com/anthropics/financial-services>
-

21) Final takeaway

If you know nothing about finance, think of this project like this:

It is a **research operating system**.

You type in an entity or portfolio. The system finds the right records. It proves what it finds with evidence. It maps the important relationships. It analyzes investment implications. It drafts a structured report. Humans improve and verify the report. Then the system keeps watching for what changes next.

That is the product.

22) Source basis for this handbook

This handbook was synthesized from: - the uploaded developer-handoff specification, - the uploaded software development plan, - prior analysis and implementation planning context from the ongoing discussion, - official public documentation for major data-source and architecture components.

The uploaded docs remain the source of truth for project scope and requirements. This handbook is a plain-English and technical interpretation intended to make execution easier.